



AI for Everyone

*Presidential Initiative for Artificial Intelligence and
Computing*

Objectives of this Course

1. The meaning behind common AI terminology, including neural networks, machine learning, deep learning, and data science
2. What AI realistically can--and cannot--do
3. How to spot opportunities to apply AI to problems in your own organization
4. What it feels like to build machine learning and data science projects
5. How to work with an AI team and build an AI strategy in your company
6. How to navigate ethical and societal discussions surrounding AI

\$13 Trillion

AI value creation by 2030

[Source: McKinsey Global Institute](#)

A lot of the value created by AI will be outside the software industry. AI will have a huge impact on all the major industries.

There are 2 types of AI

ANI

Artificial Narrow Intelligence

LOTS OF PROGRESS

ALMOST NO PROGRESS

AGI

Artificial General Intelligence

Artificial Narrow Intelligence (ANI)

These are AIs that do one thing such as:

- smart speaker
- self-driving car
- AI to do web search
- AI applications in farming or in a factory.

These types of AI are one trick ponies but when you find the appropriate trick, this can be incredibly valuable.

Artificial General Intelligence (AGI)


That is the goal to build AI.

They can do anything a human can do or maybe even be superintelligent and do even more things than any human can.



Progress in ANI vs AGI

The rapid progress in ANI has caused people to conclude that there's a lot of progress in AI, which is true. But that has caused people to falsely think that there might be a lot of progress in AGI as well which is leading to some irrational fears about evil clever robots coming over to take over humanity anytime now.



Achieving AGI Will Take Time

AGI is an exciting goal for researchers to work on, but it requires many technological breakthroughs before we get there and it may be decades or hundreds of years or even thousands of years away.

Machine Learning

Supervised Learning

If the input is an audio clip, and the AI's job is to output the text transcript, then this is speech recognition.

Input (A)
Audio

Output (B)
Text (0/1)

Application
Speech Recognition

Supervised Learning

If you want to input English and have it output a different language, Chinese, Spanish, something else, then this is machine translation.

Input (A)
English

Output (B)
Chinese

Application
Machine Translation

Supervised Learning

All the large online ad platforms have a piece of AI that inputs some information about an ad, and some information about you, and tries to predict, will you click on this ad or not?

Input (A)
Ad + User Info

Output (B)
Click? (0/1)

Application
Machine Translation

Supervised Learning

If you want to build a self-driving car, one of the key pieces of AI is the AI that takes as input an image, and some information from radar, or from other sensors, and outputs the position of other cars, so your self-driving car can avoid the other cars.



Input(A)
Image, radar info

Position of other cars

Self-Driving Car

Supervised Learning

In Manufacturing, we take as input a picture of something you've just manufactured, such as a picture of a cell phone coming off the assembly line., and you want to output, is there a scratch, or is there a dent, or some other defect on this thing you've just manufactured? This is visual inspection which is helping manufacturers to reduce or prevent defects in the things that they're making.



Input (A)
Image of a phone



Output (B)
Defects (0/1)

Visual Inspection



Supervised Learning

This set of AI called supervised learning, just learns input to output, or A to B mappings. On one hand, input to output, A to B it seems quite limiting. But when you find a right application scenario, this can be incredibly valuable.

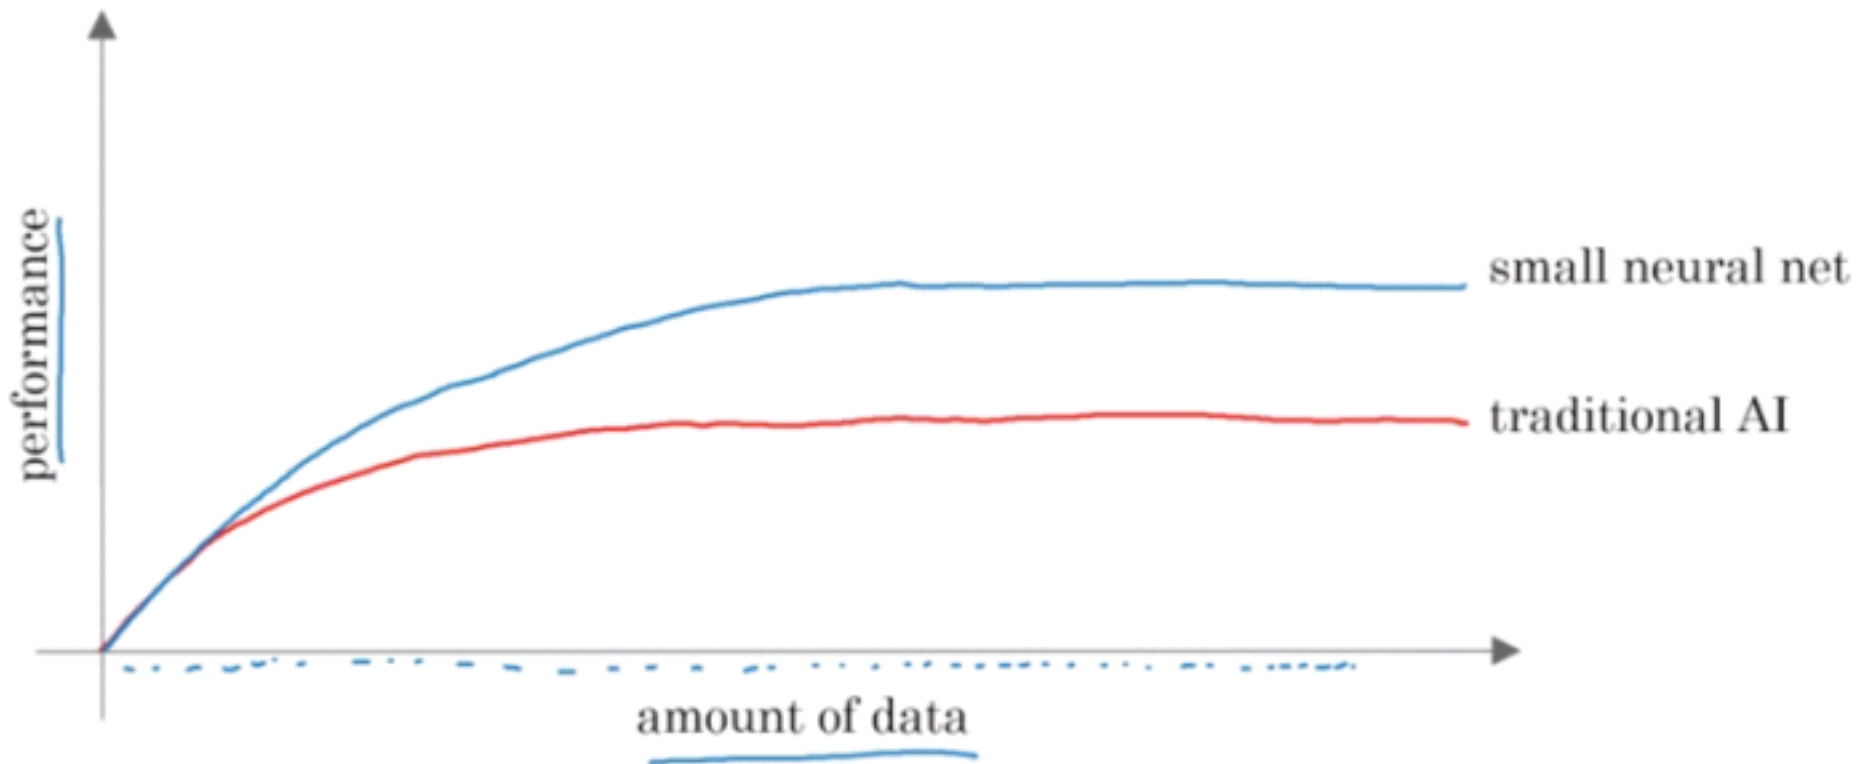


Why Now?



Why Now?


Why Now?






Why Now?

Why Now?



Why Now?





The Rise of Fast Computers

So, the rise of fast computers with specialized processors such as graphics processing units or GPUs has enabled many companies, not just giant tech companies, but many many other companies to be able to train large neural nets on a large enough amount of data in order to get very good performance and drive business value.

What is the most important
idea in AI?

Machine Learning

What is Supervised Learning?

A to B mappings

Input to Output mappings

What enables machine learning
to work so well?

What is Data

Collection of information (Facts and Figure)



A Table of Data (Dataset)

Size of House (Square Feet)	Price (\$1000)
523	115
645	150
708	210
1034	280
2290	355
2545	440
A	B

A Table of Data (Dataset)

Size of House (Square Feet)	# of Bedrooms	Price (\$1000)
523	1	115
645	1	150
708	2	210
1034	3	280
2290	4	355
2545	4	440
A		B



Data is often unique to your business

Data is often unique to your business, and this is an example of a dataset that a real estate agency might have that they tried to help price houses.

It's up to you to decide what is A and what is B, and how to choose these definitions of A and B to make it valuable for your business.

Another example

If you have a certain budget and you want to decide what is the size of house you can afford, then you might decide that the input A is how much does someone spend and B is just the size of the house in square feet, and that would be a totally different choice of A and B that tells you, given a certain budget, what's the size of the house you should be maybe looking at.

A Table of Data (Dataset)

Size of House (Square Feet)	# of Bedrooms	Price (\$1000)
523	1	115
645	1	150
708	2	210
1034	3	280
2290	4	355
2545	4	440
B		A



Acquiring data

- Manual labeling

Acquiring data

- From observing behaviors of humans

User ID	Time	Price (\$)	Purchased
4783	Jan 21 08:15.20	7.95	yes
3893	Mar 3 11:30.15	10.00	yes
8384	Jun 11 14:15.05	9.50	no
0931	Aug 2 20:30.55	12.90	yes

Acquiring data


- From observing behaviors of machines

Machine	Temperature	Pressure (psi)	Machine Fault
17987	60	7.65	N
34672	100	25.50	N
08542	140	75.50	Y
98536	165	125	Y
Input A			Input B

Acquiring data

- Download from websites / partnerships
 - Thanks to the open internet you can find so many datasets available for free online
 - Computer vision or image datasets
 - Self driving car datasets
 - Speech recognition datasets
 - Medical imaging datasets
 - Keep in mind licensing and copyright

Use and misuse of data




Give me three years to build up my IT team, we're collecting so much data.

Then after three years, I'll have this perfect dataset.

We'll do AI then.

What's wrong with this approach?

Use and misuse of data



It turns out that's a really bad strategy.

Once you've started collecting some data, go ahead and start showing it or feeding it to an AI team.

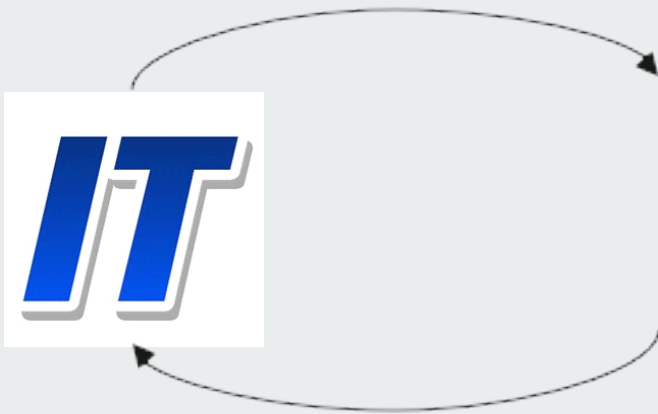
Then the AI team can give feedback to your IT team on what types of data to collect and what types of IT infrastructure to keep on building.

Example

Maybe an AI team can look at your factory data and say, "Hey. You know what? If you can collect data from this big manufacturing machine, not just once every ten minutes, but instead once every one minute, then we could do a much better job building a preventative maintenance systems for you."

Machine	Temperature	Pressure (psi)	Machine Fault
17987	60	7.65	N
34672	100	25.50	N
08542	140	75.50	Y
98536	165	125	Y
Input A			Input B


Use and misuse of data



"Hey, I have so much data. Surely, an AI team can make it valuable."

What's wrong with this statement?

Use and misuse of data




Unfortunately, this doesn't always work out.

More data is usually better than less data, but I wouldn't take it for granted that just because you have many terabytes or gigabytes of data, that an AI team can actually make that valuable.

Don't throw data at an AI team and assume it will be valuable.

Data is Messy



Not a cat

Cat

Not a cat

Cat

Data is Messy

If you have bad data, then the AI will learn inaccurate things.

Data problems:

- Incorrect labels
- Missing values

Multiple types of data

- Unstructured Data: Images, audio, text

Example

You can have incorrect labels or just incorrect data. For example, this house is probably not going to sell for \$0.1 just for one dollar.

Or, data can also have missing values such as we have here a whole bunch of unknown values.

This is structured data.

Size of House (Square Feet)	# of Bedrooms	Price (\$1000)
523	1	115
645	1	0.001
708	unknown	210
1034	3	unknown
unknown	4	355
2545	unknown	440



Machine Learning vs Data Science

Size of House (Square Feet)	# of Bedrooms	# of Bathrooms	Newly Renovated	Price (\$1000)
523	1	2	N	115
645	1	3	N	150
708	2	1	N	210
1034	3	3	Y	280
2290	4	4	N	355
2545	4	5	Y	440
A				B



Running AI System

A software that which automatically returns output B for input A.

If you have an AI system running, serving dozens or hundreds of thousands or millions of users, that's usually a machine-learning system.

Data Science

If you want to have a team analyze your dataset in order to gain insights. The output of a data science project is a set of insights that can help you make business decisions

So, a team might come up with conclusions like:

- "Hey, did you know if you have two houses of a similar size, they've a similar square footage, if the house has three bedrooms, then they cost a lot more than the house of two bedrooms, even if the square for this is the same."

Data Science

- "Did you know that newly renovated homes have a 15% premium, and this can help you make decisions such as, given a similar square footage, do you want to build a two bedroom or three bedroom size in order to maximize value? "
- "Is it worth an investment to renovate a home in the hope that the renovation increases the price you can sell a house for?"

The output of a data science project is a set of insights that can help you make business decisions, such as what type of house to build or whether to invest in renovation.



Machine Learning vs Data Science

Machine Learning

“Field of study that gives computers the ability to learn without being explicitly programmed.”

- Arthur Samuel (1959)

A machine learning project will often result in a piece of software that runs, that outputs B given A.



Formal Definition of Data Science

Data science is the science of extracting knowledge and insights from data.

So, the output of a data science project is often a slide deck, the presentation summarizes conclusions for executives to take business actions or summarizes conclusions for a product team to decide how to improve a website.



Example of ML vs DS in the online ad industry

Large platforms have AI that quickly tells them what's the ad you're most likely to click on. This is a machine learning system. It inputs information about the user and about the ad and outputs whether the user will click on the ad or not.

These systems run 24/7 and drive ad revenue for these platforms.




Example of ML vs DS in the online ad industry

If analyzing data tells you, for example, that the travel industry is not buying a lot of ads, but if you send more salespeople to sell ads to travel companies, you could convince them to use more advertising, then that would be an example of a data science project.

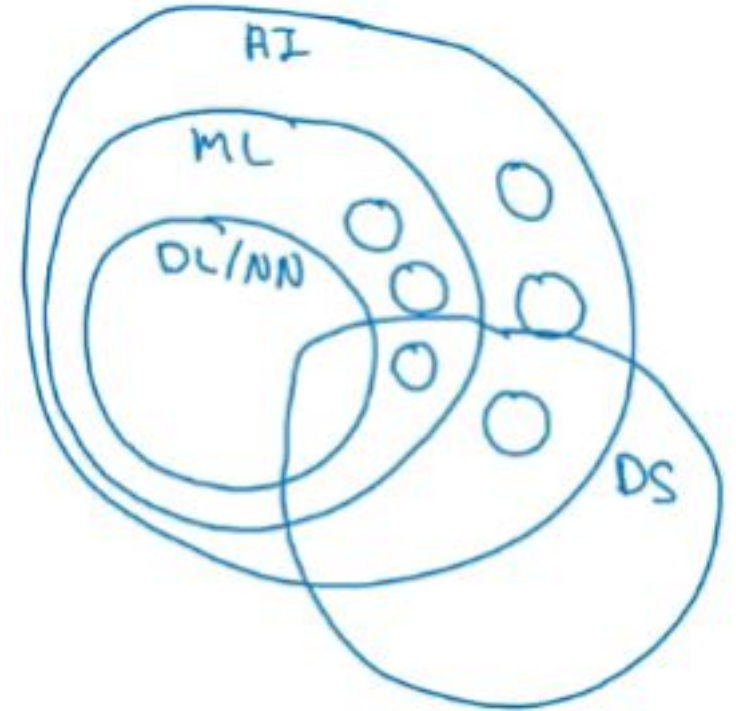
The data science conclusion results in the executives deciding to ask a sales team to spend more time reaching out to the travel industry.

Deep Learning



AI and related disciplines

- Machine Learning
- Data Science
- Deep Learning / Neural Network
- Supervised Learning
- Un supervised learning
- Reinforcement Learning





What makes a company AI company?

- Strategic data acquisition
- Unified datawarehouse
- Pervasive automation
- New roles such as MLE

AI Transformation

1. Execute pilot projects to gain momentum
2. Build an in-house AI team
3. Provide broad AI training
4. Develop an AI strategy
5. Develop internal and external communications

Deciding about a new project

- Technical diligence
 - Is it feasible project?
 - Can AI do that?
- Pretty much any thing you can do with a second of thought can be automated using supervised learning

Supervised learning tasks

Input (A)	Output (B)	Application
email	spam? (0/1)	spam filtering
audio	text transcripts	speech recognition
English	Chinese	machine translation
ad, user info	click? (0/1)	online advertising
image, radar info	position of other cars	Self-driving car
image of phone	defect? (0/1)	visual inspection

What machine learning today can and cannot do

The toy arrived two days late, so I wasn't able to give it to my niece for her birthday.

Can I return it?



"Refund request"



Oh, sorry to hear that.
I hope your niece had a good birthday.
Yes, we can help with....

Input text → Refund/Shipping/Other

A → B



Examples of what ML can and can't do?

- Identifying the intent of the customer - Possible
- Writing an emphatic response to customer's email – Not possible or difficult



Technical diligence rules

- You are learning a simple concept
- Do you have large training data

More examples

- Self driving car
 - Input is from sensors, camera
 - Output where are the other cars
- Recognizing gesture of traffic police, construction work, people
 - not possible
 - Critical application requires good accuracy

X-ray diagnosis

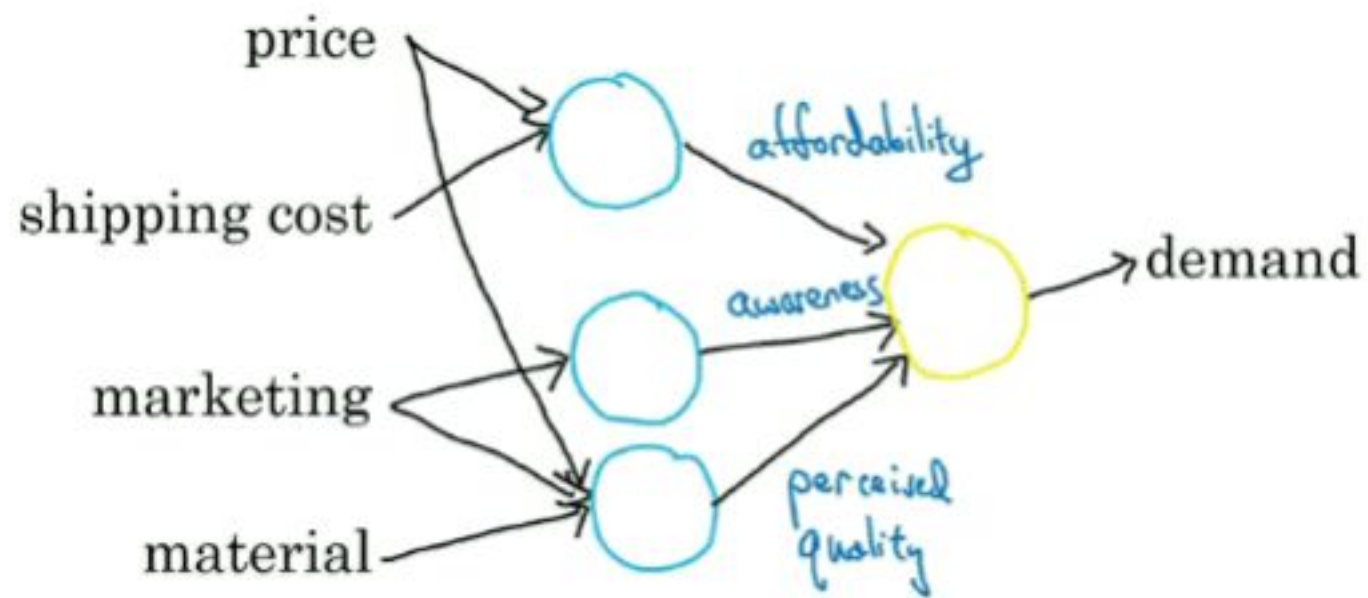
- Diagnosing a disease from X-ray images– possible
- Diagnosing a disease after reading a book

Strengths and weakness of ML

- Works when,
 - Learning a simple concept
 - Lots of data available
- Doesn't work when,
 - Learning a complex concept
 - Asked to work on new type of data such as X-ray images in different conditions and angles

Demand prediction based on price

- Price -> Demand can be modeled using a neural network using a neuron
 - (Perceptron model)
- Network of neurons (ANN)
 - Price
 - Shipping Cost
 - Marketing
 - Material



Face recognition

- Pictures comprise pixels
 - Color images and channels
- A neural network corresponds to pixels
- Earlier layers will detect edges, then lobes and then objects

Face recognition




1000

1000

1,000,000

3,000,000

A





Workflow of a Machine Learning Project

Speech Recognition



Amazon
Echo / Alexa



Google
Home



Apple
Siri



Baidu
DuerOS

Key steps of Echo / Alexa

- Collect data
 - Labelled voice
- Train model
 - Iterate many times
- Deploy the model
 - Get more data and update model

Key steps of a machine learning project

Self-driving car

1. Collect data



image → position of other cars

2. Train model

Iterate many times until
good enough



3. Deploy model

Get data back
Maintain / update model



Data Science Project Workflow

Example: Optimizing a sales funnel



Key steps of a data science project

Optimizing a sales funnel

1. Collect data

2. Analyze data

Iterate many times to get good insights

3. Suggest hypotheses/actions


Deploy changes

Re-analyze new data periodically

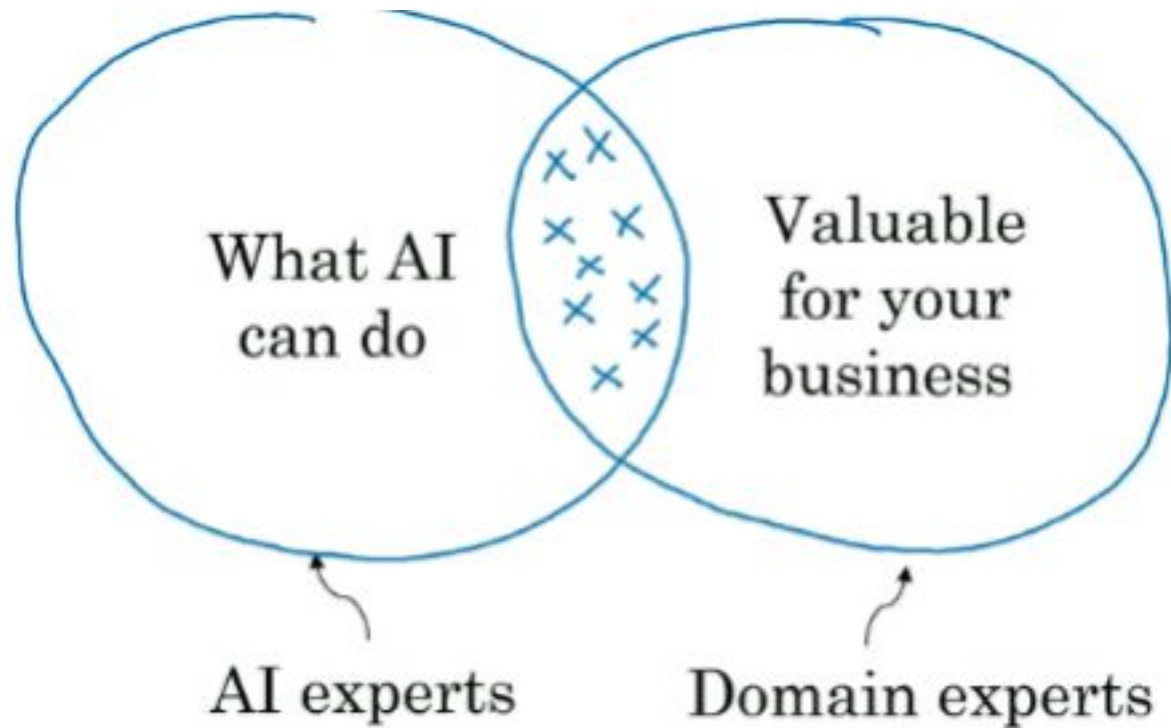
User ID	Country	Time	Webpage
2009	Spain	08:34:30 Jan 5	home.html
2897	USA	13:20:22 May 18	redmug.html
4893	Philippines	22:45:16 Jun 11	mug.html

Machine Learning changing job functions

- Sales
 - Identifying sales opportunities
 - Prioritizing
- Manufacturing line manager
 - Optimize manufacturing
 - Machine learning can spot defects
- Recruiting
 - Identify how people prefer recruitment
 - Spot good candidates

- 
- Marketing
 - Optimize website
 - A/B testing
 - Recommendation system
 - Agriculture
 - What to plant?
 - Precision agriculture

How to choose an AI project?



Brainstorming framework

- Automate task rather than job
 - Automating call center: picking phone, emails, issue refund, call routing
 - Automating radiologist: X-ray, mentoring other doctors, consulting,
- Main drivers of business value
- What are the main pain points in your business?



Is it always necessary to have big data?

- Having more data is good
- With small datasets you can make progress
- 10, 100 or 1000 data points can be a good start

Technical diligence

- Can AI system meet desired performance
- How much data is needed
- Engineering timeline

Business diligence

- Lower costs
- Increase revenue
- Launch new product or business



Ethical diligence

- Is this going to make society better?

Build Vs Buy

- ML projects can be inhouse or outsourced
- DS projects are generally inhouse
- Buy industry standard, only build specialized products

How to work with AI team

- Specify your acceptance criteria
 - 95% accuracy
 - Training, validation and Test dataset
- Don't expect 100% accuracy
 - Limitations of ML
 - Insufficient data
 - Mislabeled data
 - Ambiguous labels (human perception)

Machine Learning frameworks

Machine learning frameworks:

- TensorFlow
- PyTorch
- Keras
- MXNet
- CNTK
- Caffe
- PaddlePaddle
- Scikit-learn
- R
- Weka

Research publications:

- Arxiv

Open source repositories

- GitHub

CPU Vs GPU

CPU: Computer processor (Central Processing Unit)



GPU: Graphics Processing Unit



Cloud vs. On-premises

Edge
Deployment